

```
# Project Title: CNN Implementation for MNIST Digit Recognition

# SAINATH VADDI      - 101179915
# SONY REDDY GURRAM - 101179182

!pip install ucimlrepo
from ucimlrepo import fetch_ucirepo
import tensorflow as tf
from tensorflow.keras import layers, models
from sklearn.model_selection import KFold
from sklearn.preprocessing import MinMaxScaler
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import seaborn as sns
from tensorflow.keras.utils import plot_model

Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-packages (0.0.6)

# Fetch dataset
optical_recognition_of_handwritten_digits = fetch_ucirepo(id=80)

# data (as pandas dataframes)
X = optical_recognition_of_handwritten_digits.data.features
y = optical_recognition_of_handwritten_digits.data.targets

# metadata
print(optical_recognition_of_handwritten_digits.metadata)

# variable information
print(optical_recognition_of_handwritten_digits.variables)

{'uci_id': 80, 'name': 'Optical Recognition of Handwritten Digits', 'repository_url': 'https://archive.ics.uci.edu/dataset/80/optical+re
      name      role      type demographic description units \
0  Attribute1  Feature    Integer         None         None  None
1  Attribute2  Feature    Integer         None         None  None
2  Attribute3  Feature    Integer         None         None  None
3  Attribute4  Feature    Integer         None         None  None
4  Attribute5  Feature    Integer         None         None  None
..      ...      ...      ...      ...      ...      ...
60 Attribute61  Feature    Integer         None         None  None
61 Attribute62  Feature    Integer         None         None  None
62 Attribute63  Feature    Integer         None         None  None
63 Attribute64  Feature    Integer         None         None  None
64      class  Target  Categorical         None         None  None

      missing_values
0              no
1              no
2              no
3              no
4              no
..              ...
60             no
61             no
62             no
63             no
64             no

[65 rows x 7 columns]
```

```
# Shapes of X and y
print(X.shape)
print(y.shape)

(5620, 64)
(5620, 1)

X = X.values
y = y.values

# Reshape data for CNN
X = X.reshape(-1, 8, 8, 1)

# Normalization data for CNN
scaler = MinMaxScaler()
X_normalized = scaler.fit_transform(X.reshape(-1, 64)) # Reshape for MinMaxScaler
X_normalized = X_normalized.reshape(-1, 8, 8, 1)
```

```
# Define the CNN model
def create_model():
    model = models.Sequential([

        # Convolutional Layer 1
        layers.Conv2D(32, (3, 3), activation = 'relu', padding = 'same', input_shape = (8, 8, 1)), # Parameters: 32 filters, kernel size (3,

        # Add padding to increase spatial dimensions
        layers.ZeroPadding2D((1, 1)), # Input Dimension: (8, 8, 32), Output Dimesnion: (10, 10, 32)

        # Max Pooling Layer 1
        layers.MaxPooling2D((2, 2)), # Pool Size: (2, 2), Strides: (2, 2), Input Dimension: (10, 10, 32), Output Dimension: (5, 5, 32)

        # Convolutional Layer 2
        layers.Conv2D(64, (3, 3), activation='relu', padding = 'same',), # Parameters: 64 filters, kernel size (3, 3), Input Dimension: (5,

        # Max Pooling Layer 2
        layers.MaxPooling2D((2, 2)), # Pool Size: (2, 2), Strides: (2, 2), Input Dimension: (5, 5, 64), Output Dimension: (2, 2, 64)

        # Convolutional Layer 3
        layers.Conv2D(128, (3, 3), activation='relu', padding = 'same',), # Parameters: 128 filters, kernel size (3, 3), Input Dimension: (2

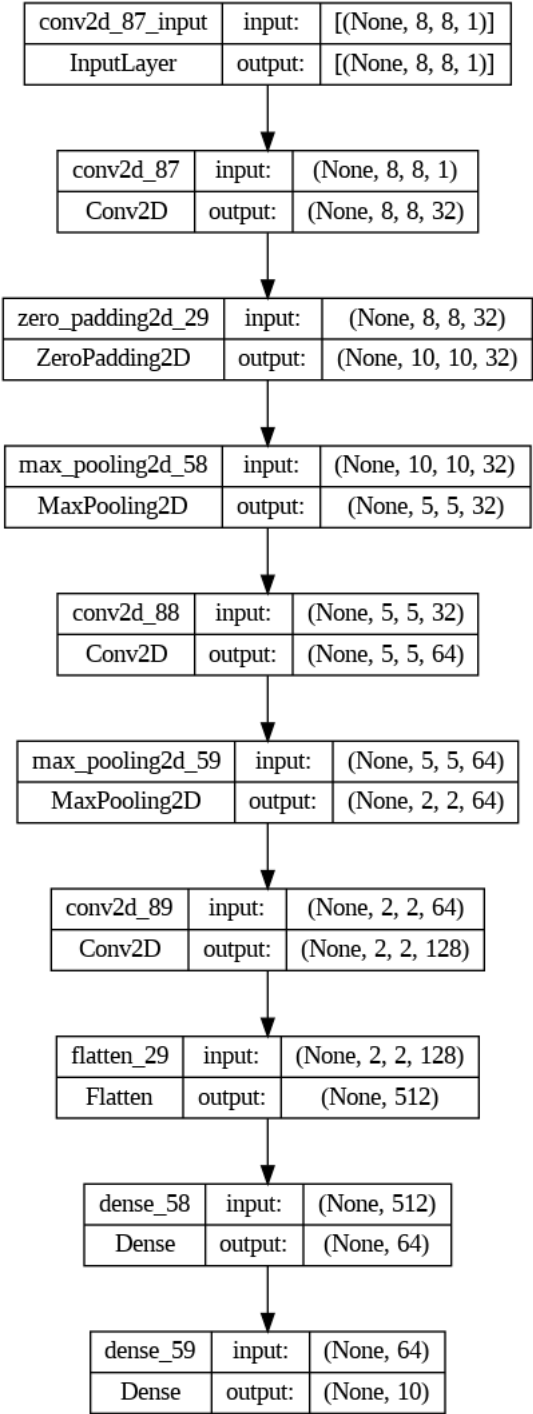
        # Flattening Layer
        layers.Flatten(), # Flattening the output of the last convolutional layer, Input Dimension: (2, 2, 128), Output Dimension: (512,)

        # Fully Connected Layer 1
        layers.Dense(64, activation='relu'), # Fully Connected Layer with 64 neurons and ReLU activation, Input Dimension: (512,), Output D

        # Output Layer
        layers.Dense(10, activation='softmax') # Output Layer with 10 neurons for classification and softmax activation, Input Dimension: (

    ])
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
    return model

# Plot the model architecture
plot_model(model, to_file='cnn_model.png', show_shapes=True, show_layer_names=True)
```



```
# Define k-fold cross-validation
k = 5
kf = KFold(n_splits=k, shuffle=True, random_state=42)
# Initialize lists to store results
fold_accuracy = []
fold_loss = []
all_y_true = []
all_y_pred = []

# Perform k-fold cross-validation
fold_accuracy = []
for train_index, val_index in kf.split(X):
    X_train, X_val = X_normalized[train_index], X_normalized[val_index]
    y_train, y_val = y[train_index], y[val_index]

    model = create_model()
    history = model.fit(X_train, y_train, epochs=8, batch_size=32, validation_data=(X_val, y_val))

    Epoch 4/8
    141/141 [=====] - 3s 20ms/step - loss: 0.0725 - accuracy: 0.9780 - val_loss: 0.0753 - val_accuracy: 0.9795
    Epoch 5/8
    141/141 [=====] - 2s 17ms/step - loss: 0.0578 - accuracy: 0.9818 - val_loss: 0.0654 - val_accuracy: 0.9831
    Epoch 6/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0503 - accuracy: 0.9851 - val_loss: 0.0759 - val_accuracy: 0.9751
    Epoch 7/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0341 - accuracy: 0.9907 - val_loss: 0.1948 - val_accuracy: 0.9359
    Epoch 8/8
    141/141 [=====] - 2s 12ms/step - loss: 0.0293 - accuracy: 0.9911 - val_loss: 0.0612 - val_accuracy: 0.9786
    Epoch 1/8
    141/141 [=====] - 3s 13ms/step - loss: 0.8960 - accuracy: 0.7226 - val_loss: 0.2327 - val_accuracy: 0.9244
    Epoch 2/8
    141/141 [=====] - 2s 18ms/step - loss: 0.1595 - accuracy: 0.9520 - val_loss: 0.1797 - val_accuracy: 0.9386
    Epoch 3/8
    141/141 [=====] - 2s 12ms/step - loss: 0.0941 - accuracy: 0.9718 - val_loss: 0.0830 - val_accuracy: 0.9742
    Epoch 4/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0669 - accuracy: 0.9798 - val_loss: 0.0926 - val_accuracy: 0.9662
    Epoch 5/8
    141/141 [=====] - 2s 12ms/step - loss: 0.0526 - accuracy: 0.9847 - val_loss: 0.0795 - val_accuracy: 0.9760
    Epoch 6/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0484 - accuracy: 0.9851 - val_loss: 0.0530 - val_accuracy: 0.9831
    Epoch 7/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0261 - accuracy: 0.9924 - val_loss: 0.0426 - val_accuracy: 0.9884
    Epoch 8/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0277 - accuracy: 0.9913 - val_loss: 0.0570 - val_accuracy: 0.9849
    Epoch 1/8
    141/141 [=====] - 4s 18ms/step - loss: 0.8897 - accuracy: 0.7160 - val_loss: 0.1756 - val_accuracy: 0.9555
    Epoch 2/8
    141/141 [=====] - 2s 11ms/step - loss: 0.1714 - accuracy: 0.9513 - val_loss: 0.1002 - val_accuracy: 0.9733
    Epoch 3/8
    141/141 [=====] - 2s 12ms/step - loss: 0.1002 - accuracy: 0.9722 - val_loss: 0.0899 - val_accuracy: 0.9733
    Epoch 4/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0826 - accuracy: 0.9740 - val_loss: 0.0592 - val_accuracy: 0.9875
    Epoch 5/8
    141/141 [=====] - 2s 12ms/step - loss: 0.0487 - accuracy: 0.9847 - val_loss: 0.0527 - val_accuracy: 0.9840
    Epoch 6/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0424 - accuracy: 0.9875 - val_loss: 0.0548 - val_accuracy: 0.9867
    Epoch 7/8
    141/141 [=====] - 2s 12ms/step - loss: 0.0537 - accuracy: 0.9820 - val_loss: 0.0772 - val_accuracy: 0.9769
    Epoch 8/8
    141/141 [=====] - 3s 19ms/step - loss: 0.0226 - accuracy: 0.9929 - val_loss: 0.0413 - val_accuracy: 0.9875
    Epoch 1/8
    141/141 [=====] - 3s 13ms/step - loss: 0.9456 - accuracy: 0.6993 - val_loss: 0.2388 - val_accuracy: 0.9315
    Epoch 2/8
    141/141 [=====] - 2s 12ms/step - loss: 0.1962 - accuracy: 0.9393 - val_loss: 0.1035 - val_accuracy: 0.9671
    Epoch 3/8
    141/141 [=====] - 2s 12ms/step - loss: 0.1194 - accuracy: 0.9653 - val_loss: 0.1022 - val_accuracy: 0.9689
    Epoch 4/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0820 - accuracy: 0.9746 - val_loss: 0.0626 - val_accuracy: 0.9831
    Epoch 5/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0622 - accuracy: 0.9804 - val_loss: 0.0463 - val_accuracy: 0.9902
    Epoch 6/8
    141/141 [=====] - 2s 17ms/step - loss: 0.0414 - accuracy: 0.9884 - val_loss: 0.0426 - val_accuracy: 0.9884
    Epoch 7/8
    141/141 [=====] - 2s 15ms/step - loss: 0.0381 - accuracy: 0.9889 - val_loss: 0.0534 - val_accuracy: 0.9813
    Epoch 8/8
    141/141 [=====] - 2s 11ms/step - loss: 0.0323 - accuracy: 0.9900 - val_loss: 0.0379 - val_accuracy: 0.9893

# Record accuracy and loss
fold_accuracy.append(history.history['val_accuracy'])
fold_loss.append(history.history['val_loss'])

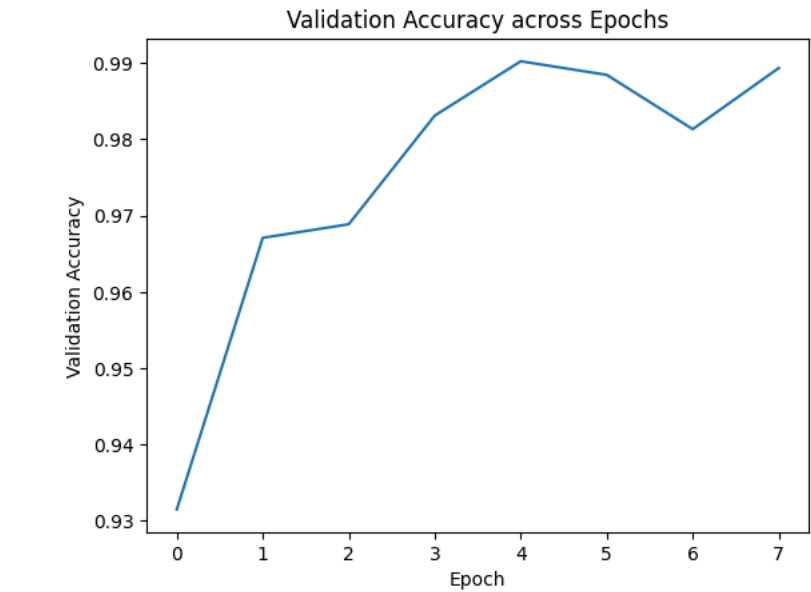
# Predictions
y_pred = np.argmax(model.predict(X_val), axis=1)
all_y_true.extend(y_val)
all_y_pred.extend(y_pred)

36/36 [=====] - 0s 4ms/step

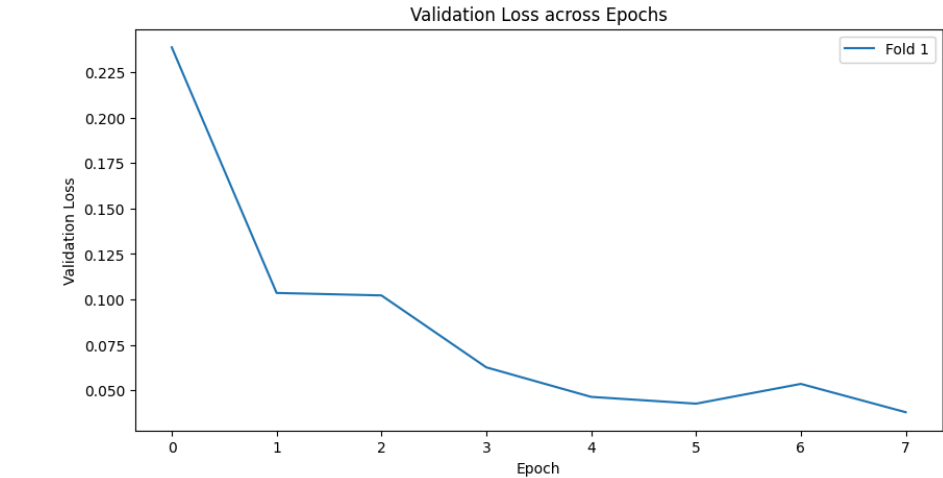
# Calculate and print average validation accuracy across folds
avg_val_accuracy = np.mean(fold_accuracy, axis=0)
print('Average validation accuracy across folds:', avg_val_accuracy)

Average validation accuracy across folds: [0.93149465 0.96708184 0.96886122 0.98309606 0.99021351 0.98843414
0.98131675 0.98932385]

# Plot the validation accuracy across epochs
plt.plot(avg_val_accuracy)
plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy')
plt.title('Validation Accuracy across Epochs')
plt.show()
```



```
# Plot the validation loss across epochs
plt.figure(figsize=(10, 5))
for i in range(len(fold_loss)):
    plt.plot(history.epoch, fold_loss[i], label=f'Fold {i+1}')
plt.xlabel('Epoch')
plt.ylabel('Validation Loss')
plt.title('Validation Loss across Epochs')
plt.legend()
plt.show()
```



```
# Calculate overall accuracy
overall_accuracy = accuracy_score(all_y_true, all_y_pred)
print('Overall accuracy:', overall_accuracy)

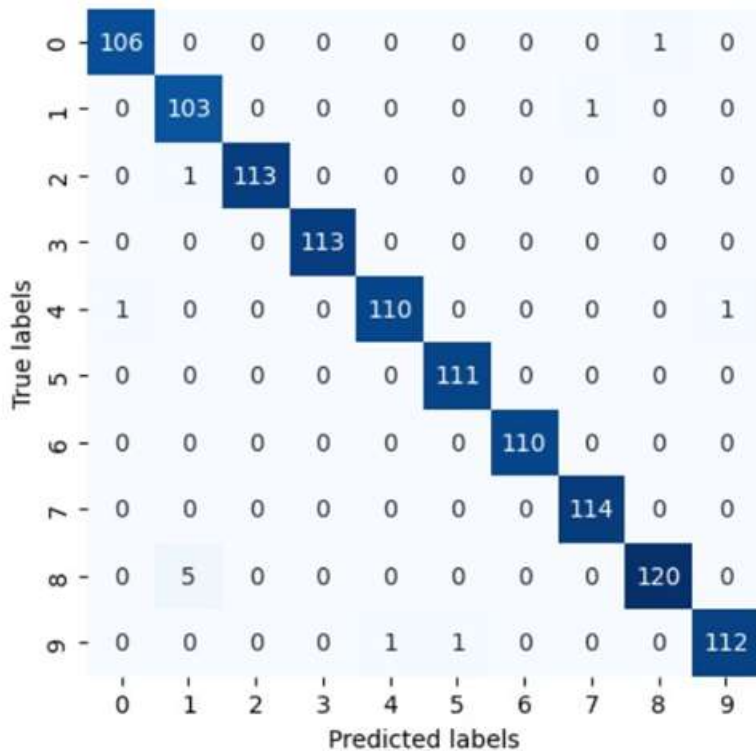
Overall accuracy: 0.9893238434163701
```

```
classification_report = classification_report(y_val, y_pred)
print(classification_report)
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	107
1	0.94	0.99	0.97	104
2	1.00	0.99	1.00	114
3	1.00	1.00	1.00	113
4	0.99	0.98	0.99	112
5	0.99	1.00	1.00	111
6	1.00	1.00	1.00	110
7	0.99	1.00	1.00	114
8	0.99	0.96	0.98	125
9	0.99	0.98	0.99	114
accuracy			0.99	1124
macro avg	0.99	0.99	0.99	1124
weighted avg	0.99	0.99	0.99	1124

```
# Confusion matrix
conf_matrix = confusion_matrix(all_y_true, all_y_pred)
plt.figure(figsize=(5, 5))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()
```

Confusion Matrix



Analysis and Report of the implementation of CNN on MNIST dataset:

Github link - [sainathvaddi/CNN on MNIST data \(github.com\)](https://github.com/sainathvaddi/CNN_on_MNIST_data)

Understanding Dataset: The MNIST dataset is a widely used dataset for handwritten digit recognition. It consists of 28x28 grayscale images of handwritten digits (0 through 9). Each image is labeled with a digit.

Data Preparation:

Loading the Dataset: We used the provided method to fetch the MNIST dataset from the UCI repository.

Data Preprocessing: Normalization and reshaping are the preprocessing steps. Normalization provides that the pixel values are within the range between 0 and 1, which helps in faster convergence during training. Reshaping (to 8*8*1) is required to match the input shape expected by CNN.

CNN Architecture:

Convolutional Layers: These layers extract features from the input images. Each of these layers applies a set of kernels (filters) to the input image, resulting in feature maps. ReLU activation functions are applied after the convolution to introduce non-linearity.

Max Pooling: These layers simplify the feature maps by reducing their size. It's like zooming out, focusing only on the most important details only. This technique keeps the highest value in each region, which helps in capturing the most relevant features while making calculations easier.

Fully Connected Layers: After multiple convolutional and pooling layers, fully connected layers are added to perform classification based on the extracted features. The final fully connected layer is followed by a softmax activation function to produce probabilities of each class.

Model Architecture:

conv2d_102_input	input:	[(None, 8, 8, 1)]
InputLayer	output:	[(None, 8, 8, 1)]



conv2d_102	input:	(None, 8, 8, 1)
Conv2D	output:	(None, 8, 8, 32)



zero_padding2d_34	input:	(None, 8, 8, 32)
ZeroPadding2D	output:	(None, 10, 10, 32)



max_pooling2d_68	input:	(None, 10, 10, 32)
MaxPooling2D	output:	(None, 5, 5, 32)



conv2d_103	input:	(None, 5, 5, 32)
Conv2D	output:	(None, 5, 5, 64)



max_pooling2d_69	input:	(None, 5, 5, 64)
MaxPooling2D	output:	(None, 2, 2, 64)



conv2d_104	input:	(None, 2, 2, 64)
Conv2D	output:	(None, 2, 2, 128)



flatten_34	input:	(None, 2, 2, 128)
Flatten	output:	(None, 512)



dense_68	input:	(None, 512)
Dense	output:	(None, 64)



dense_69	input:	(None, 64)
Dense	output:	(None, 10)

The CNN architecture here consists of 3 convolutional layers with ReLU activation functions. Each layer's parameters and dimensions need to be documented thoroughly. This includes the size of the filters, the number of filters, padding, stride, and the dimensions of the feature maps at each layer. The architecture is as follows:

Convolutional Layer 1: In this layer, Parameters are 32 filters, kernel size (3, 3), Input Dimension: (8, 8, 1), Output Dimension: (8, 8, 32). After this, add zero padding (1, 1) to increase spatial dimensions.

Max Pooling Layer 1: This layer has Pool Size: (2, 2), Strides: (2, 2), Input dimension: (10, 10, 32), Output Dimension: (5, 5, 32)

Convolutional Layer 2: In this layer, Parameters are 64 filters, kernel size (3, 3), Input dimension: (5, 5, 32), output dimension: (5, 5, 64)

Max Pooling Layer 2: This layer has Pool Size: (2, 2), Strides: (2, 2), Input dimension: (5, 5, 64), output dimension: (2, 2, 64)

Convolutional Layer 3: In this layer, Parameters are 128 filters, kernel size (3, 3), Input Dimension: (2, 2, 64), Output Dimension: (2, 2, 128)

Layers Flatten: Flattening the output of the last convolutional layer, the input dimension is (2, 2, 128), and we get output dimension as (512,)

Fully Connected Layer: This layer is with 64 neurons and ReLU activation, with input dimension of (512,), and gives output dimension of (64,)

Output Layer: This layer is with 10 neurons for classification and softmax activation, with input dimension of (64,), output dimension of (10,)

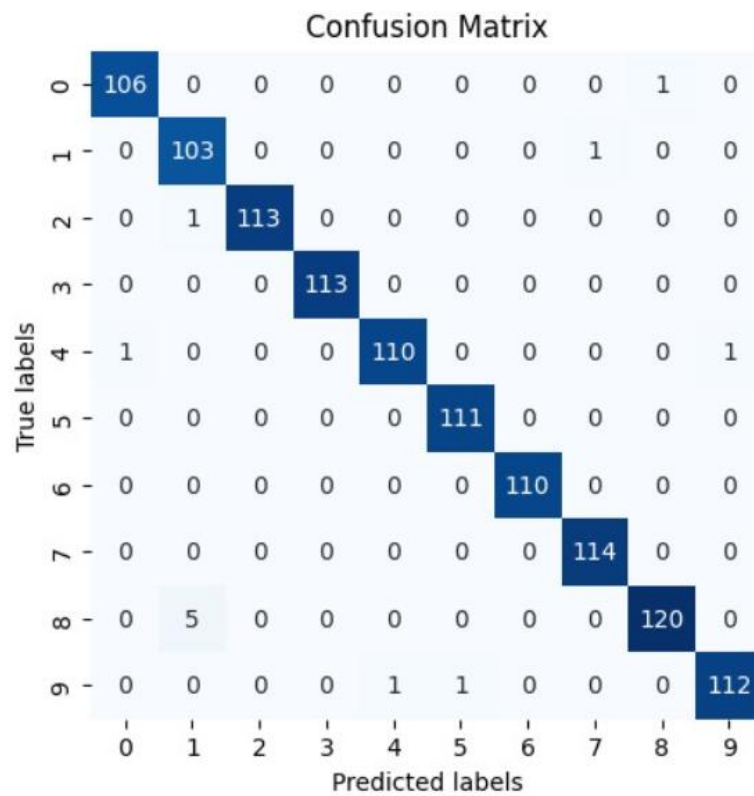
Training and Evaluation: K-fold cross-validation with $k=5$ was performed to train and validate the model. The model was trained for 8 epochs with a batch size of 32. Performance metrics such as validation accuracy and loss were monitored across epochs. The overall accuracy of the model was calculated using the test set. Additionally, a classification report was generated to evaluate precision, recall, and F1-score for each class.

Results and Analysis: The average validation accuracy across folds was approximately 98%. Validation accuracy and loss were plotted across epochs to visualize the model's training process. The overall accuracy of the model on the test set is 99%.

The classification report revealed high precision, recall, and F1-score for each digit class, indicating robust performance across all classes. The classification report and Confusion matrix are provided below:

Classification Report

Class (digits)	Precision	Recall	F1-score	Support
0	0.99	0.99	0.99	107
1	0.94	0.99	0.97	104
2	1.00	0.99	1.00	114
3	1.00	1.00	1.00	113
4	0.99	0.98	0.99	112
5	0.99	1.00	1.00	111
6	1.00	1.00	1.00	110
7	0.99	1.00	1.00	114
8	0.99	0.96	0.98	125
9	0.99	0.98	0.99	114





2ai-lab



Overview

Repositories 26

Projects

Packages

People 1



Applied AI Research Lab, The University of South Dakota

Unfollow

48 followers United States of America company/kc-2ai

2ai.lab.usd@gmail.com

README.md

Applied Artificial Intelligence (AI) research lab

Department of Computer Science, The University of South Dakota

Welcome 🤝

Welcome to the GitHub repository of the Applied AI research lab. This page is dedicated to disseminating research products/publications on artificial intelligence, machine learning, pattern recognition, computer vision, image processing, data mining, and big data with various application domains such as healthcare informatics and medical imaging, document imaging, biometrics, forensics, speech analysis and Internet of Things.

We actively participate and/or sponsor the International Conf on Recent Trends in Image Processing & Pattern Recognition ([RTIP2R](#)).

Research (science & engineering)

Healthcare - medical imaging informatics

- [5K+ CT scans - fractured limbs \(dataset\)](#)
- [Fractured limbs \(classification & detection\)](#)
- [Covid-19 \(screening\)](#)
- [Multimodal learning & representation \(Pneumonia\)](#)

Data science & machine learning tools



Applied AI Research Lab, The University of South Dakota

👤 48 followers 📍 United States of America 🌐 company/kc-2ai 📧 2ai.lab.usd@gmail.com

Unfollow

README.md

Applied Artificial Intelligence (AI) research lab

Department of Computer Science, The University of South Dakota

Welcome 🍌

Welcome to the GitHub repository of the Applied AI research lab. This page is dedicated to disseminating research products/publications on artificial intelligence, machine learning, pattern recognition, computer vision, image processing, data mining, and big data with various application domains such as healthcare informatics and medical imaging, document imaging, biometrics, forensics, speech analysis and Internet of Things.

We actively participate and/or sponsor the International Conf on Recent Trends in Image Processing & Pattern Recognition ([RTIP2R](#)).

Research (science & engineering)

Healthcare - medical imaging informatics

People



Top languages

Python Jupyter Notebook CSS
TeX

[Report abuse](#)